

Unsupervised Learning

Clustering, Dimensionality Reduction, Data Preparation

kwekha.rostam.zhyar@nik.uni-obuda.hu

06

Unsupervised Learning



What is the difference?

- Supervised learning: The model learns from labeled data (x_i, y_i)
- Unsupervised learning: The model searches for structures and hidden patterns in unlabeled data $x_i \in X$

Examples:

- Supervised learning: Image classification (e.g., cat or dog?)
- Unsupervised learning: Automatic image grouping based on similarity without predefined labels

Definition:

- Dataset: $D = \{x_n\}_{n=1}^N$, where $x_n \in \mathbb{R}^d$
- Data originates from an unknown distribution
- Goal: Learn useful properties from the unknown distribution

Hidden (latent) variables:

- Unsupervised learning often assumes that unobserved (hidden) variables explain the structure of the data

Clustering

Definition:

- Discover hidden categories within the data
- Assign each data point a hidden label z_n

Example:

- Grouping documents by topic (politics, sports, economics)
- Automatically grouping images without knowing their content beforehand

Important: Clustering is different from classification: classification involves labeled data, while clustering groups data in an unsupervised manner.

Task:

- Given a dataset $D = \{x_n\}_{n=1}^N$, where $x_n \in \mathbb{R}^d$
- The goal is to partition the data into K clusters
- The clusters are not predefined; they are discovered by the algorithm

Key Idea:

- Group elements based on their proximity
- Assign a prototype point to each cluster: μ_k , which represents the cluster center

Optimization Problem in K-Means

Formal Definition:

- Assign a hidden variable to each data point: which cluster does it belong to?
- Objective function:

$$\min_{\{z_n\}, \{\mu_k\}} \sum_{n=1}^N \sum_{k=1}^K z_{kn} d(x_n, \mu_k)$$

- Where:
 - $z_{kn} = 1$, if x_n belongs to the k -th cluster; otherwise, 0.
 - $d(x, \mu) = ||x - \mu||^2$ is the Euclidean distance.

Core Idea:

- Assign data based on minimal distance
- Continuously update cluster centers μ_k

Steps of the K-Means Algorithm

Expectation-Maximization (EM) Structure:

- 1. Initialization: Randomly select K cluster centers $\{\mu_k^{\text{old}}\}$
- 2. E-Step (Assignment):

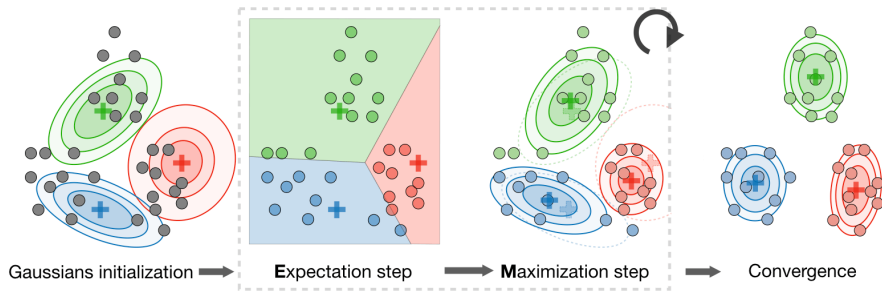
$$z_{kn}^{\text{new}} = \begin{cases} 1, & \text{if } k = \arg \min_j d(x_n, \mu_j^{\text{old}}) \\ 0, & \text{otherwise} \end{cases}$$

- 3. M-Step (Update):

$$\mu_k^{\text{new}} = \frac{\sum_{n=1}^N z_{kn}^{\text{new}} x_n}{\sum_{n=1}^N z_{kn}^{\text{new}}}$$

- 4. Convergence check: The algorithm runs until the new centers no longer change significantly.

EM Structure




1. Figure: The steps of the EM algorithm structure: initialization, assignment (E), update (M), convergence check ¹

¹<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-unsupervised-learning>

K-Means Clustering Step by Step

Example: $K=3$ clusters, convergence step by step



The figure illustrates the steps of K-Means clustering:

- **Initialization:** Random centroids
- **Assignment:** Assign data points to the nearest cluster
- **Update:** Compute new cluster centroids
- After several iterations, final clusters are formed

How to Choose the Value of K ?

The problem:

- Too small $K \rightarrow$ Too few clusters, poor grouping
- Too large $K \rightarrow$ Every data point could have its own cluster, making no sense


Determining K :

- Fine-tuning K until a predefined heuristic is met
- Hierarchical clustering (dendrogram): tree-structured representation

Main Idea:

- Construct a nested hierarchy of clusters
- The result is a dendrogram showing the order of cluster merging
- Two main approaches:
 - The **agglomerative** approach builds from the bottom up—each sample starts as its own cluster and then merges
 - The **divisive** method works oppositely, starting from the top and gradually splitting into smaller clusters

Clustered Iris data set
(the labels give the true flower species)



K-Means vs GMM

- K-Means performs hard assignments—each point belongs to a single cluster
- GMM applies probabilistic clustering—each point has a probability of belonging to a cluster
- K-Means assumes spherical clusters, while GMM can model elliptical clusters

GMM:

- Assumes data comes from multiple Gaussian distributions
- Each cluster is modeled as a Gaussian distribution
- Uses the Expectation-Maximization (EM) algorithm for parameter optimization

Mathematical Model of GMM

Gaussian Distribution

$$\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

GMM: The data comes from K Gaussian distributions:

$$p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x | \mu_k, \Sigma_k)$$

Explanation:

- π_k : Cluster weight (probability of belonging to a given cluster)
- μ_k : Cluster center (mean)
- Σ_k : Covariance matrix (variance and correlation)

Mathematical Model of GMM

$$p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x \mid \mu_k, \Sigma_k)$$

- $p(x)$ – The probability that a data point x comes from the entire mixture model.
- $\mathcal{N}(x \mid \mu_k, \Sigma_k)$ – The density function of the k -th Gaussian distribution, where:
 - μ_k : The mean of the distribution
 - Σ_k : Covariance matrix (shape, direction, variance)
- π_k – The weight of the k -th cluster (proportion of data belonging to it)

Σ_k determines the cluster:

- Size (larger variance = wider cluster),
- Shape (round or elongated),
- Orientation (can be rotated).

Expectation-Maximization Algorithm

- E-Step (Expectation): Compute probabilities of a point belonging to each cluster:

$$\gamma_{ik} = \frac{\pi_k \mathcal{N}(x_i | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_i | \mu_j, \Sigma_j)}$$

- $\gamma_{ik} \in [0, 1]$, and $\sum_{k=1}^K \gamma_{ik} = 1$

Expectation-Maximization Algorithm

M-step: The parameters are updated based on the weighted average of γ_{ik} :

$$N_k = \sum_{i=1}^N \gamma_{ik}$$

$$\mu_k^{\text{new}} = \frac{1}{N_k} \sum_{i=1}^N \gamma_{ik} x_i$$

$$\Sigma_k^{\text{new}} = \frac{1}{N_k} \sum_{i=1}^N \gamma_{ik} (x_i - \mu_k)(x_i - \mu_k)^T$$

$$\pi_k^{\text{new}} = \frac{N_k}{N}$$

Comparison of K-means and GMM

Similarities

- Both algorithms follow the **EM structure**:
 - **E-step**: assignment based on current parameters
 - **M-step**: update of cluster centroids
- Iterative convergence: each step reduces the loss function
- Assign data points to clusters

Differences

- **K-means**:
 - Hard assignment (x_i belongs to exactly one cluster)
 - Learns only cluster centroids (μ_k)
 - Assumes clusters are spherical
- **GMM**:
 - Soft assignment (probabilistic)
 - Learns cluster weights (π_k), centroids (μ_k), and covariances (Σ_k)
 - More flexible cluster shapes (e.g., ellipses)

DBSCAN (Density-Based Spatial Clustering of Applications with Noise)

- K-Means and GMM assume specific cluster shapes.
- DBSCAN does not require the number of clusters to be predefined.
- Handles clusters of different shapes and noise data.

Key concepts:

- **Core sample:** A point with at least minPoints neighbors within radius ϵ .
- **Border sample:** A neighbor of a core point, but not a core point itself.
- **Outlier:** A point that does not belong to any cluster.

Algorithm steps:

- ① Check the core point condition for each point: how many neighbors does it have within radius ϵ ?
- ② If a point is a core point, assign all reachable points to the same cluster.
- ③ If a point is a border point, assign it to the nearest cluster.
- ④ If a point is neither a core nor a border point, mark it as noise.

Visualization:

<https://www.naftaliharris.com/blog/visualizing-dbscan-clustering/>

Advantages:

- Does not require a predefined number of clusters.
- Can handle arbitrarily shaped clusters, unlike K-Means, which assumes spherical clusters.
- Automatically detects noise data.

Disadvantages:

- Sensitive to parameters: Properly setting ϵ is crucial.
- Difficult to apply in high-dimensional spaces because distances become less informative.

Dimensionality Reduction

Definition:

- Goal: Reduce the number of features while preserving the most important information
- A commonly used technique: Principal Component Analysis (PCA)

Why is it useful?

- Helps visualize data in lower dimensions
- Reduces noise and removes redundant data
- Improves computational efficiency on large datasets

Example:

- In an image database, retaining key colors and shapes while removing noisy details

Principal Component Analysis (PCA)

Goal: Reduce the dimensionality of data while preserving the maximum variance.

- PCA finds new axes (principal components) along which the variance of the data is maximized.
- The method projects the data onto these axes.
- Suitable for:
 - visualization (2D/3D),
 - noise reduction,
 - preprocessing before supervised learning.

- ① **Normalize the data:** Scale each attribute so that its mean is 0 and standard deviation is 1.

$$x_j^{(i)} = \frac{x_j^{(i)} - \mu_j}{\sigma_j}$$

- ② **Compute the covariance matrix:**

$$\Sigma = \frac{1}{m} \sum_{i=1}^m x^{(i)} (x^{(i)})^T$$

- ③ **Compute principal components:** Eigenvectors and eigenvalues of Σ .
- ④ **Project:** Project data onto the axes corresponding to the largest eigenvalues.


Attributes are often correlated \rightarrow redundant information is present.

Example: YouTube Content Creators

- Each data point represents a content creator.
- $x_1^{(i)}$: Technical skills (editing, audio, video quality).
- $x_2^{(i)}$: Enthusiasm/motivation (consistency, energy, authenticity).
- These characteristics are often correlated:
 - Motivated individuals tend to learn technical skills.
 - Good technical background enhances enjoyment and engagement.
- Goal: How can we describe these on a single scale?

PCA: Motivation

- The two features (technical knowledge and motivation) are highly correlated.
- We can assume that the data is distributed along a **diagonal axis**—this is the \vec{u}_1 direction.
- This direction may represent the combined characteristic of "content creator karma."
- Question: **How can we automatically find this axis?**



Step 1: Normalizing the Data

- For each j -th attribute:

$$x_j^{(i)} = \frac{x_j^{(i)} - \mu_j}{\sigma_j}$$


where


$$\mu_j = \frac{1}{n} \sum_{i=1}^n x_j^{(i)}, \quad \sigma_j^2 = \frac{1}{n} \sum_{i=1}^n (x_j^{(i)} - \mu_j)^2$$


- Result: each attribute will have a mean of 0 and a standard deviation of 1.
- Why is this important?
 - Different measurement units can distort the analysis (e.g., speed vs. number of seats).
 - Scaling ensures that all features contribute equally.
- When can this step be skipped?
 - If we know that the data is already zero-centered (e.g., audio signals).
 - If all attributes are on the same scale (e.g., pixel values in an image).

Step 2: Determining the Principal Direction

- After normalizing the data, we seek the direction \mathbf{u} that:
 - Is a **unit vector** ($\|\mathbf{u}\| = 1$),
 - and maximizes the **variance of the projected values**.
- Intuition:
 - The data contains a certain amount of variation (information).
 - The goal is to find the direction with the greatest variance—this represents the most significant "direction of change."
- This direction becomes the first principal component.








Selecting the Best Principal Component


On a Sample Dataset:

- The dataset is already normalized.
- We examine two different axis directions (u) onto which we project the data.


1. Normalized Data



2. Good Direction



3. Poor Direction



Summary:

- The first direction (u_1) maximizes the variance of the projected data.
- The second direction (u_2) does not align well with the data distribution.
- Goal: Automatically select the best u .

Principal Component Selection

Given an u unit vector and a data point x , the projection length is: $x^T u$

The goal is to maximize the variance of all projections:

$$\frac{1}{n} \sum_{i=1}^n (x^{(i)T} u)^2 =$$

$$\frac{1}{n} \sum_{i=1}^n u^T x^{(i)} x^{(i)T} u = u^T \left(\frac{1}{n} \sum_{i=1}^n x^{(i)} x^{(i)T} \right) u$$

The expression in parentheses is the covariance matrix:

$$\Sigma = \frac{1}{n} \sum_{i=1}^n x^{(i)} x^{(i)T}$$

This is an eigenvector-eigenvalue problem:

$$\Sigma u = \lambda u$$

Solution: u is the eigenvector corresponding to the largest eigenvalue of Σ .

Step 3: Covariance Matrix and Eigenvectors

- Compute the covariance matrix:

$$\Sigma = \frac{1}{m} \sum_{i=1}^m x^{(i)} x^{(i)T}$$

- This is a symmetric matrix, where:
 - Eigenvalues \rightarrow Measure of variance
 - Eigenvectors \rightarrow Principal component directions
- The first principal component (u_1) direction maximizes the variance along the data.

PCA Summary and Dimensionality Reduction

- If the goal is to find a 1-dimensional subspace for approximating the data, the optimal choice is:

u_1 = Eigenvector corresponding to the largest eigenvalue

- If we wish to project data onto a k -dimensional subspace ($k < d$), we choose the first k eigenvectors:

$$u_1, u_2, \dots, u_k$$

- The u_i vectors form a new, orthogonal basis.

Representing Data in the New Basis:

- The new k -dimensional representation of a data point $x^{(i)}$ is:

$$y^{(i)} = \begin{bmatrix} u_1^T x^{(i)} \\ u_2^T x^{(i)} \\ \vdots \\ u_k^T x^{(i)} \end{bmatrix} \in \mathbb{R}^k$$

- The original $x^{(i)}$ was d -dimensional, now reduced to k -dimensions.
- The vectors u_1, \dots, u_k are called the **first k principal components** of the data.

1. Dimensionality Reduction and Visualization

- If $k = 2$ or $k = 3$, the data can be visually represented.

2. Preprocessing for Machine Learning Models

- Dimensionality reduction decreases computational costs.
- Helps prevent overfitting: lower-dimensional inputs lead to simpler hypothesis classes.

3. Noise Reduction and Pattern Recognition

- Face recognition: PCA projects 100×100 pixel faces into a lower-dimensional space (*eigenfaces* method).
- Removes noise caused by small lighting variations and imaging differences.
- Distance-based similarity measurement in the reduced-dimensional space: successful face comparison.

What is anomaly detection?

- Algorithms that identify data points that significantly deviate from "normal" patterns.
- Typical applications: fraud detection, fault detection, network security.

Two popular solutions:

- **One-Class SVM:** defines a boundary around the data and considers points outside it as anomalies.
- **Isolation Forest:** isolates unusual points using random decision trees; efficient for large datasets.

Data Preparation

Data sets in tutorials



Data sets in the wild



<https://nc233.com/tag/feature-engineering/>

- Every ML project relies on a dataset.
- Possibility: open datasets (e.g., ImageNet: 14M+ labeled images, 20k+ categories).
- If no suitable dataset is available:
 - **Self-collection**: manually, via crowdsourcing, or automatically.
 - **Data from user behavior**.
 - **Transfer learning**: general model + fine-tuning on own data.

Selecting Suitable Data

- Is the data relevant to the task?
- Does it contain the required **output values**?
- Data size matters:
 - Is **more data collection** necessary?
 - Can some data be discarded to speed up processing?

How Much Data is Enough?

- Analyzing similar projects: estimation by analogy.
- **Learning curve** analysis: does performance still improve with more data?

Rule of thumb:

- Difficult tasks require millions of examples.
- For classification: at least hundreds or thousands of examples per class.
- At least 10x as many data points as input attributes.
- More data is required for nonlinear models than linear ones.
- Less data may be sufficient if regularization is used.


Data Cleaning: Potential Issues

- Are there errors in data entry?
- Missing fields? Need to determine how to handle them.
- Inconsistent textual data: "Apple" vs "AAPL" vs "Apple Inc."

Reliable data cleaning and preprocessing pipelines are the foundation of good models.

Data Augmentation

- Useful when data is scarce – especially for images, text, and speech.
- For images, for example:
 - Rotation, flipping.
 - Translation, cropping, scaling.
 - Adding noise, adjusting brightness.
- The label remains unchanged – making the model more robust.



<https://pyimagesearch.com/2019/07/08/keras-imagedatagenerator-and-data-augmentation/>

- Example: among 10 million transactions, only 1000 are fraudulent.
- A "classifying all as valid" model achieves 99.99% accuracy – but is not useful.
- Solutions:
 - **Undersampling**: reducing the majority class.
 - **Oversampling**: duplicating the minority class.
 - **Weighted loss function**: higher penalty for misclassification.
 - **Boosting**: minority decisions may suffice.
 - **SMOTE, ADASYN**: generating synthetic examples.

- Outliers are distant data points (e.g., a \$316 price when others are around \$30).
- Linear models are sensitive to them.
- **Solutions:**
 - Logarithmic transformation: e.g., $20 \rightarrow 1.3$, $316 \rightarrow 2.5$
 - Decision trees, random forests are more robust to outliers.

What is Feature Engineering?

- The transformation or augmentation of attributes to make learning easier for the model.
- Goal: create **relevant**, **informative**, and **processable** features.
- Follows after correcting erroneous values.

Example:

- *Wait time* → discrete categories: 0–10, 10–30, 30–60, >60 minutes.
- *Date* → logical fields such as “weekend?” or “holiday?”

Typical Transformations and Encodings

- **Quantization (binning):** continuous values \rightarrow categories
- **Normalization:** standard deviation = 1 (particularly important for algorithms like k-NN)
- **One-hot encoding:** categories \rightarrow binary fields
 - E.g., "weather = sunny/cloudy/rainy" \rightarrow 3 separate attributes

Important: Some models (e.g., neural networks) can only handle numerical inputs.

Example of Quantization (Binning)

Problem: The "Wait Time" attribute is given in minutes:

Wait Time (minutes) = 4, 8, 12, 35, 50, 65

Goal: Transform into categories:

- Short (0–10 minutes)
- Medium (10–30 minutes)
- Long (30+ minutes)

Result:

4 \rightarrow Short, 12 \rightarrow Medium, 50 \rightarrow Long

Example of One-Hot Encoding

Attribute: Weather = sunny, cloudy, rainy

One-Hot Encoding:

Original Value	sunny	cloudy	rainy
sunny	1	0	0
cloudy	0	1	0
rainy	0	0	1

Note: Models can interpret the categories in numeric form this way.

Example: Real Estate Price Estimation

- Base feature: floor area → not enough.
- Additional features:
 - Number of rooms, bedrooms, bathrooms
 - Renovations, year built, condition
 - Heating, air conditioning, garden size, orientation

Environmental Factors:

- Postal code? School district? Average test scores?
- “Good neighborhood” is not always numerical – but can be treated as a feature.

The Importance of Good Features

- The quality of features is critical to the success of the model.
- A weak model can perform well with good features.
- Even a strong model cannot learn from poor representations.

Pedro Domingos:

“At the end of the day, some machine learning projects succeed and some fail. What makes the difference? Easily the most important factor is the features used.”

Thank you for your attention!

Sources:

- Osvaldo Simeone, A Brief Introduction to Machine Learning for Engineers.
- <https://scikit-learn.org/stable/modules/clustering.html>
- Andrew Ng, CS229 Lecture notes, Stanford University.
- Russell, Norvig, Artificial Intelligence: A Modern Approach, 4th Edition.